

Compiler Optimisation Coursework

s1208506

The University of Edinburgh

February 14, 2018

Abstract

This report details the use of an iterative algorithm in attempting to find an optimal selection of compiler flags for GCC targetting twelve specific benchmark programs. A set of weights for each flag were continually updated to improve the execution speed over time. Results for each benchmark are presented as well as comparing each flag set over the cumulative time. For each benchmark, many flag sets were found which improved the result over -O3 and two sets were able to beat -O3 for all twelve benchmarks.

1 Introduction

When compiling a program using the GCC compiler, several preset levels of optimisation are available: -O0, -O1, -O2, -O3. These implement a selection of *compiler flags* which have been found to be almost always beneficial in improving the execution time of the resulting program. The highest level, -O3, is more experimental and can sometimes decrease performance over -O2. This is due to the difficulty in creating a single set of compiler flags which are effective for any target program.

This report describes a method for iteratively selecting additional compiler flags beyond -O3, in an attempt to both find an optimal set of flags for each target program as well as the best set across all of the test programs. A set of varied benchmark programs were provided as the target programs for this investigation.

2 Evaluation Methodology

2.1 Initial Flag Selection

In order to reduce the problem space, a subset of all possible GCC compiler optimisation flags was selected

from the official documentation page [1]. The compiler options from -O1 and -O2 have been extensively tested so were always included in the tests. The flags in the -O3 set are considered more experimental, so for each experiment set they were set active by default but the corresponding -fno flags were included in the option pool. A collection of other flags which were not in any of the -O1, -O2 or -O3 sets were also included in this option pool. Optimisation parameters (**param**) were not considered in this investigation, so all flags used their default values. In total, a pool of 65 flags was chosen.

2.2 Flag Weighting Method

In an effort to iteratively improve the result in subsequent runs, a flag weighting system was devised. For each benchmark program, the execution time for -O3 was measured and used as the target time. All individual flag options were then run alone¹ for each benchmark and the resulting execution time compared against -O3. The results for these tests were used as the initial *weight* values assigned to each flag where the weight value was equal to the integer number of milliseconds faster than -O3². A distinct set of weights was maintained for each benchmark in addition to an *overall* set which compared the cumulative time across all 12 benchmark programs.

The hypothesis was that flags which performed well should be included more often in test sets. Therefore the next test set should be drawn from the total pool of flags based upon probabilities proportional to the current weight assigned to each flag. After a test run, the weights for all flags which were selected should be updated depending on if the execution time was greater or less than that of -O3. To make this system more stable, a *damping factor* of 0.25 was introduced. Each flag was then increased or decreased by the execution time difference (compared to -O3) multiplied by 0.25.

¹In addition to -O3

²Weights can be negative if the flag was slower than base -O3

For each iteration, either the *overall* set of weights or a specific benchmark is targetted and used as the probability distribution for drawing flags from the pool. The number of flags drawn was chosen to be a random integer between 2 and 50. Once a set was chosen, regardless of the target benchmark, it was run through all 12 benchmarks as per the experiment specification. Thus each iteration would update all 12 individual flag weight sets as well as the overall set.

2.3 Experiments

The experimental procedure followed the schedule:

1. Run each benchmark for -00, -01, -02 and -03. (x4)
2. Run each flag alone as -03 [flag]. (x65)
3. Run 50 iterations of the selection algorithm targetting the *overall* pool. (x50)
4. Run 20 iterations targetting each benchmark in turn. (x240)
5. Run 50 iterations of the selection algorithm targetting the *overall* pool. (x50)

Therefore 409 iterations were scheduled to run. However, as some flags were incompatible with each other, any selection which failed to compile or resulted in an execution output which was inconsistent with the result from -00 was discarded. This left 344 successful runs.

3 Experimental Setup

3.1 Platform

The experiments were all conducted on an Apple Macbook Air laptop. In an effort to reduce any processor *noise*, all other programs were closed and the laptop was disconnected from the internet. Table 1 details the exact specification of the hardware used.

The operating system used was macOS High Sierra, version 10.13.2. All scripting was done using Python 3.6.4 and the version of GCC used was 7.3.0.

3.2 Timing

It was important to accurately measure the execution time of each compiled benchmark program. Each program was run multiple times in order to average the run times and reduce the effects of noise. The

Processor name	Intel Core i7 (4650U)
Clock Speed	1.7 Ghz
Number of cores	2
L2 Cache per core	256 KB
L3 Cache	4 MB
Main Memory	8 GB DDR3 @ 1600MHz

Table 1: Hardware specification used in the following tests

mean and standard deviation were calculated but the primary method for ranking the run times was the median. For measuring the execution time of the target benchmarks, the results were expected to not be normally distributed as there is a bound on the lower run time (fastest possible execution) but the upper limit is unbounded [2]. At least ten iterations were performed to comply with the provided specification. However, additional iterations were run in order to ensure that the 95% confidence interval in the median was less than 3% of the median wide.

The 95% lower confidence limit was found as the i^{th} value of the sorted list of execution times where i was found as:

$$i = \frac{n}{2} - \frac{1.96\sqrt{n}}{2}$$

Similarly, the 95% upper confidence limit was found as the j^{th} value of the sorted list where j was found as:

$$j = 1 + \frac{n}{2} + \frac{1.96\sqrt{n}}{2}$$

When updating the weights after each benchmark, if the median execution time for the test set fell within the lower and upper confidence bounds for -03, then the result was considered equivalent to -03 and the weights remained unchanged.

To measure the raw timing values for each run the UNIX `time` command was used. This returns values for *real*, *user* and *sys* time for the program under test. For the purposes of this experiment, it was decided that *real* time should be used as it represents raw “wall-clock” time. An alternative option was to use *user* + *sys* time which would return the amount of time spent by the CPU actually running the process (*user*) and spent in the kernel within the process (*sys*). Results from some initial experiments showed that $real \approx user + sys$. If a program was able to use more than one core then it would be possible for *user* + *sys* time to exceed *real* time. As it was unknown whether any parallelisation was present in the benchmark programs or could be added via certain compiler

flags, *real* time was chosen, although this could potentially suffer greater effects from external CPU activity.

4 Results

For each benchmark, the results will be presented in a table. The median values were the key measurement for analysis. Statistics for the four base optimisation levels will be given, as well as data for the fastest flag set. The average time from all iterative flag sets is also given. A full list of flags in each fastest flag set are also listed in the appendix.

4.1 401.bzip2

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	7166	68	7168	7094	7289
-01	3541	23	3542	3503	3568
-02	3384	26	3379	3360	3433
-03	3373	37	3375	3334	3433
Fastest	3157	31	3151	3146	3183
Average	-	-	3351	-	-

SD: Standard Deviation; **Med.:** Median

CI_l : 95% Confidence Interval in the median lower limit;

CI_u : 95% Confidence Interval in the median upper limit;

Table 2: Execution time results for 401.bzip2 in ms

Results are summarised in Table 2. The fastest set completed in 3151 ms which was 6.6% faster than -03, which completed in 3375 ms. These flags are given in the appendix as flag set 1.

The average execution time was 3351 ms, 0.7% faster than -03.

4.2 429.mcf

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	3643	59	3644	3612	3694
-01	2873	253	2780	2752	2829
-02	2339	33	2341	2311	2376
-03	2350	27	2351	2324	2391
Fastest	2282	27	2285	2248	2313
Average	-	-	2354	-	-

Table 3: Execution time results for 429.mcf in ms

Results are summarised in Table 3. The fastest set completed in 2285 ms which was 2.8% faster than -03,

which completed in 2351 ms. -02 was also faster than -03 with a time of 2341 ms, but the fastest set was still 2.4% faster. This fastest flag set is given in the appendix as flag set 2.

The average execution time was 2354 ms, 0.13% slower than -03.

4.3 433.milc

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	24670	136	24636	24524	24858
-01	15321	233	15400	15159	15486
-02	14807	331	14870	14582	15017
-03	14544	337	14577	14325	14748
Fastest	13039	218	12963	12916	13237
Average	-	-	14106	-	-

Table 4: Execution time results for 433.milc in ms

Results are summarised in Table 4. The fastest set completed in 12963 ms which was 11.1% faster than -03, which completed in 14577 ms. This flag set is given in the appendix as set 3. The average execution time was 14106 ms, 3.2% faster than -03.

4.4 445.gobmk

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	25867	729	25603	25576	25988
-01	15145	65	15148	15071	15252
-02	14658	70	14680	14559	14739
-03	14422	53	14423	14350	14499
Fastest	14058	52	14050	14004	14176
Average	-	-	14525	-	-

Table 5: Execution time results for 445.gobmk in ms

Results are summarised in Table 5. The fastest set completed in 14050 ms which was 2.6% faster than -03, which completed in 14423 ms. This flag set is listed in the appendix as set 4.

The average execution time was 14525 ms, 0.71% slower than -03.

4.5 456.hmmmer

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	7801	45	7799	7754	7879
-01	2377	46	2357	2352	2420
-02	2099	41	2091	2070	2123
-03	2140	55	2121	2110	2164
Fastest	2086	50	2071	2057	2118
Average	-	-	2151	-	-

Table 6: Execution time results for 456.hmmmer in ms

Results are summarised in Table 6. The fastest set completed in 2071 ms which was 4.9% faster than -03, which completed in 2121 ms. -02 was also faster than -03, with a time of 2091 ms, but the fastest set was still 0.96% faster. This flag set is listed in the appendix as set 5.

The average execution time was 2151 ms, 1.4% slower than -03.

4.6 458.sjeng

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	5415	50	5399	5385	5458
-01	3258	28	3258	3226	3298
-02	3201	24	3202	3190	3243
-03	3236	20	3234	3218	3269
Fastest	3099	28	3088	3072	3143
Average	-	-	3206	-	-

Table 7: Execution time results for 458.sjeng in ms

Results are summarised in Table 7. The fastest set completed in 3088 ms which was 4.8% faster than -03, which completed in 3243 ms. -02 was also faster than -03 with a time of 3202 ms, but the fastest set was still 3.6% faster. This set of flags is given in the appendix as flag set 6.

The average execution time was 3206 ms, 1.1% faster than -03.

4.7 462.libquantum

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	3475	28	3485	3451	3509
-01	1207	20	1204	1191	1227
-02	986	14	986	972	1000
-03	984	15	982	973	998
Fastest	835	12	832	827	847
Average	-	-	1181	-	-

Table 8: Execution time results for 462.libquantum in ms

Results are summarised in Table 8. The fastest set completed in 832 ms which was 15.3% faster than -03, which completed in 982 ms. This set of flags is given in the appendix as set 7.

The average execution time was 1181 ms, 20% slower than -03. Indeed, several flag sets were in fact slower than -00.

4.8 464.h264ref

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	31771	587	31532	31440	31668
-01	10050	183	9994	9962	10055
-02	9968	83	9948	9876	10130
-03	9210	57	9213	9136	9314
Fastest	8840	54	8860	8770	8903
Average	-	-	9265	-	-

Table 9: Execution time results for 464.h264ref in ms

Results are summarised in Table 9. The fastest set completed in 8860 ms which was 3.8% faster than -03, which completed in 9213 ms. This flag set is given in the appendix as set 8.

The average execution time was 9265 ms, 0.56% slower than -03.

4.9 470.lbm

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	3137	37	3134	3106	3172
-01	2277	23	2279	2248	2314
-02	2279	25	2284	2255	2312
-03	2270	21	2270	2240	2305
Fastest	2256	22	2246	2242	2288
Average	-	-	2286	-	-

Table 10: Execution time results for 470.lbm in ms

Results are summarised in Table 10. The fastest set completed in 2246 ms which was 1.1% faster than -03, which completed in 2270 ms. This flag set is given in the appendix as set 9.

The average execution time was 2286 ms, 0.70% slower than -03.

4.10 h263dec

This benchmark could not compile due to issues in the code. It was omitted from the investigation.

4.11 h263enc

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	3770	42	3765	3739	3836
-01	1932	29	1922	1912	1966
-02	1902	21	1890	1884	1929
-03	1811	18	1810	1797	1829
Fastest	1752	21	1745	1736	1786
Average	-	-	1831	-	-

Table 11: Execution time results for h263enc in ms

Results are summarised in Table 11. The fastest set completed in 1745 ms which was 3.6% faster than -03, which completed in 1810 ms. These flags are given in flag set 10 in the appendix.

The average execution time was 1831 ms, 1.2% slower than -03.

4.12 mpeg2dec

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	313	14	310	307	315
-01	150	10	147	146	150
-02	138	12	135	133	137
-03	126	8	125	123	126
Fastest	120	12	116	115	118
Average	-	-	124	-	-

Table 12: Execution time results for mpeg2dec in ms

Results are summarised in Table 12. The fastest set completed in 116 ms which was 7.2% faster than -03, which completed in 125 ms. These flags are given in flag set 11 in the appendix.

The average execution time was 124 ms, 0.8% faster than -03.

4.13 mpeg2enc

Flag set	Mean	SD	Med.	CI_l	CI_u
-00	8097	206	8025	7965	8128
-01	2291	21	2289	2266	2328
-02	2278	29	2272	2255	2316
-03	2231	31	2233	2206	2271
Fastest	2194	26	2179	2172	2224
Average	-	-	2251	-	-

Table 13: Execution time results for mpeg2enc in ms

Results are summarised in Table 13. The fastest set completed in 2179 ms which was 2.4% faster than -03, which completed in 2233 ms. These flags are given in the appendix as flag set 12.

The average execution time was 2251 ms, 0.8% slower than -03.

4.14 Overall

Flag set	Mean	Med.
-00	125146	124518
-01	60436	60336
-02	58056	58094
-03	56713	56728
Fastest	54757	54731
Average	-	56655

Table 14: Execution time results for the cumulative execution time of all 12 benchmarks in ms

Results are summarised in Table 14. The fastest set completed in 54731 ms which was 3.5% faster than -03, which completed in 56728 ms. These flags are given in the appendix as flag set 13.

This is the “fastest set of flags across all programs” but was only faster than -03 in 10 of the 12 benchmarks.

The average execution time was 56655 ms, 0.13% faster than -03.

There were two flag sets which beat -03 in all 12 benchmarks. These were -03 **-fcx-limited-range**, which had a cumulative time of 56230 ms and another set, given as flag set 14 of the appendix, which had a cumulative time of 54865 ms.

A summary of the relative performance of the fastest and average flag sets against -03 is given in Table 15.

Benchmark	-02	-03	Fastest	Average
libquantum	986	982	-15.3%	+20%
milc	14870	14577	-11.1%	-3.2%
mpeg2dec	135	125	-7.2%	-0.80%
bzip2	3379	3375	-6.6%	+0.70%
hmmer	2091	2121	-4.9%	+1.4%
sjeng	3202	3234	-4.8%	-1.1%
h264ref	9948	9213	-3.8%	+0.56%
h263enc	1890	1810	-3.6%	+1.2%
mcf	2341	2311	-2.8%	+0.13%
gobmk	14680	14423	-2.6%	+0.71%
mpeg2enc	2272	2233	-2.4%	+0.80%
lbm	2284	2270	-1.1%	+0.70%
<i>overall</i>	58094	56728	-3.5%	-0.13%

Table 15: Summary of the fastest and average flag sets vs. -03

5 Analysis and Discussion

5.1 Observations from the results

In all benchmarks, the iterative weight-based algorithm was successful in finding multiple flag sets which produced a faster execution time of the benchmark over -03. It was interesting to note that of the 344 flag sets tested, only two were consistently faster than -03 for all 12 benchmarks. This highlights the difficulty in finding one *best* set of compiler flags for all possible programs.

For most benchmarks, the average result of all flag sets was close to the performance of -03. This was to be expected, as they all start from the baseline settings of -03. An interesting anomaly was in the *libquantum* benchmark where the average was 20% worse. There were some flag sets which produced disastrous results, even slower than -00. However, each flag that was included in some of the worst sets, with execution times around 4900 ms, was also present in some of the best sets with execution times around 850 ms. The raw timing data was examined to ensure that these results were not consecutive as this may have indicated a period of external CPU noise. These 4000+ ms runs were evenly spread throughout which suggests this was not the issue. Instead, it implies that it was the combinations of certain flags which produced these issues.

The level of success for the best performing flag sets was also drastically different for each benchmark. Improvements of over 10% were achieved for the *libquantum* and *milc* benchmarks while a speedup of only 1.1% was found for the *lbm* benchmark. It is notable that *libquantum* had both the best improvement with its fastest set and the worst performance on average. Compare this to *lbm* which featured a minor change in both cases. This suggests, at least for the chosen flag pool in this experiment, that some programs are more invariant to compilation techniques while others have greater scope for performance alterations.

Similar variability can be seen in the speed-up from -00 to -03 as shown in Table 16. However, there is no strong correlation between the benchmarks which benefit the most from the -03 optimisations and those which benefit the most from the custom optimisations.

Benchmark	-O3 vs. -O0	Fastest vs. -O3
libquantum	-72%	-15.3%
milc	-41%	-11.1%
mpeg2dec	-60%	-7.2%
bzip2	-35%	-6.6%
hmmer	-73%	-4.9%
sjeng	-40%	-4.8%
h264ref	-71%	-3.8%
h263enc	-52%	-3.6%
mcf	-35%	-2.8%
gobmk	-44%	-2.6%
mpeg2enc	-72%	-2.4%
lbm	-28%	-1.1%
<i>overall</i>	-54%	-3.5%

Table 16: Compilation benefits comparison

5.2 The Weight-based Algorithm

The use of the weight system allowed for the relative performance of each flag for each benchmark to be assessed. Although this was not perfect, as a great flag could be pushed down the rankings due to frequently being drawn along with poor flags. Consistently high performing flags were `-flto` and `-fprofile-use` while consistently poor flags were `-fno-inline-functions` and `-fno-inline`. Even with these however, there were certain benchmarks where they had the opposite effect.

Across all flag weight sets, the highest value achieved was `-flto` on the *milc* benchmark with a weight of 29934. The lowest value achieved was `-ftree-parallelize-loops=16` on the *libquantum* benchmark with a weight of -21199.

To further improve the results, some additional work could be done to the search algorithm. A major flaw in the weight-based system is that a potentially good flag could, by chance, be included in a set with one or more highly detrimental flags early in the process. If this set then resulted in a very negative result, then the potentially strong flag will have its weight reduced to an unrepresentative value and it would then be unlikely to be drawn again. More iterations targeting each benchmark would also be beneficial in further improving the results.

Another extension would be in extending the weight system to the flag set size parameter. Instead of just randomly selecting a set size each iteration, the sizes could also be weighted to investigate if smaller or larger set sizes were more likely to result in an effective set. From the results, this was also a benchmark specific property. For example, the majority of

the best performing benchmarks for *mcf* were a single flag whereas for *libquantum*, most had over 40.

6 Summary and Conclusions

This investigation has proven that using iterative compilation techniques is effective at reducing the execution time for various benchmark programs. A significant performance increase over the standard `-O3` flag setting is possible when targetting a single program. However, difficulties in finding one set of compiler flags which are effective for any arbitrary program were highlighted.

The specific flags included, and the number of flags in the fastest set, were significantly variable for each of the benchmarks investigated. This leads to the conclusion that the method of iterative compilation for a specific target program is essential in achieving an optimal result.

While analysing the results, it became clear that there were many complexities involved with selecting compiler optimisation flags, even for a single target program. Flags could produce good results alone but when used in combination, produce highly negative results and vice versa. It was also observed that some programs were more invariant to improvements from the compiler than others.

If optimum performance is sought for a target program then an iterative compilation system would be a valuable exercise.

References

- [1] Free Software Foundation, Inc. “3.10 Options That Control Optimization,” *Using the GNU Compiler Collection (GCC)* <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>
- [2] Torsten Hoefler, “How many measurements do you need to report a performance number?” <https://htor.inf.ethz.ch/blog/index.php/2016/04/14/how-many-measurements-do-you-need-to-report-a-performance-number/>

Appendix: Full flag sets

The full flag sets mentioned in the report are listed here.

Flag set	Benchmark	Flags
1	bzip2	-O3 -fsplit-loops -fbranch-probabilities -fsched-stalled-insns -fweb -fselective-scheduling -fno-predictive-commoning -floop-unroll-and-jam -fstdarg-opt -fkeep-inline-functions -funroll-loops -fno-zero-initialized-in-bss -fprefetch-loop-arrays -fno-defer-pop -floop-nest-optimize -fno-function-cse -fmerge-all-constants -flimit-function-alignment -fgcse-las -fvpt -fno-inline-functions -fsched-critical-path-heuristic -fpeel-loops -fconserve-stack -fsched-spec-load-dangerous -fipa-pta -fsched-spec-insn-heuristic -funroll-all-loops -fvariable-expansion-in-unroller -fselective-scheduling2 -fsched-group-heuristic -fsched-rank-heuristic -fmodulo-sched -fsection-anchors -fisolate-erroneous-paths-attribute -fivopts -funswitch-loops -fno-ipa-cp-clone -ftree-vectorize -freschedule-modulo-scheduled-loops -ftree-loop-im -fsched-spec-load
2	mcf	-O3 -fsection-anchors
3	milc	-O3 -fselective-scheduling -funroll-loops -fprofile-use -fsched-last-insn-heuristic -fstdarg-opt -fsched-rank-heuristic -flimit-function-alignment -fselective-scheduling2 -fmodulo-sched -floop-nest-optimize -fsched-spec-load-dangerous -fipa-pta -fgcse-after-reload -fprofile-correction -fkeep-static-functions -fprefetch-loop-arrays -flto -fbtr-bb-exclusive -funswitch-loops -fkeep-inline-functions -floop-unroll-and-jam
4	gobmk	-O3 -ftree-vectorize -fisolate-erroneous-paths-attribute -flto -flimit-function-alignment -fcx-limited-range
5	hammer	-O3 -fisolate-erroneous-paths-attribute -fno-ira-share-spill-slots -fsched-stalled-insns-dep -fno-zero-initialized-in-bss -flto -fcx-limited-range -fprofile-use -fselective-scheduling2 -fstdarg-opt -fipa-pta -fgcse-sm -fsched-dep-count-heuristic -fsplit-loops -fsched-group-heuristic -frename-registers -freschedule-modulo-scheduled-loops
6	sjeng	-O3 -fselective-scheduling -fno-ira-share-save-slots -flto -flive-range-shrinkage -fsched-stalled-insns -fsched-last-insn-heuristic -fweb -fdevirtualize-speculatively -fvpt -fgcse-sm -fno-function-cse -floop-parallelize-all -fprofile-use -fgcse-after-reload -fivopts -fstdarg-opt -flimit-function-alignment -fno-ipa-cp-clone -fprefetch-loop-arrays -funroll-all-loops -fconserve-stack -fkeep-inline-functions -fsched2-use-superblocks -freschedule-modulo-scheduled-loops -fsched-spec-insn-heuristic -fkeep-static-functions -fvariable-expansion-in-unroller -fsched-pressure -fno-zero-initialized-in-bss -fsched-spec-load

Flag set	Benchmark	Flags
7	libquantum	-O3 -fgcse-after-reload -floop-unroll-and-jam -fkeep-inline-functions -fbtr-bb-exclusive -fbranch-target-load-optimize -fivopts -fsched-spec-load-dangerous -funroll-all-loops -fpeel-loops -fsched2-use-superblocks -fprefetch-loop-arrays -flto -fweb -fselective-scheduling -fmerge-all-constants -fipa-pta -fno-ira-share-spill-slots -flimit-function-alignment -fno-function-cse -fgcse-sm -fsched-critical-path-heuristic -floop-nest-optimize -frename-registers -fbranch-probabilities -fsched-pressure -fcx-limited-range -fsched-spec-load -fsched-spec-insn-heuristic -ftree-loop-ivcanon -funroll-loops -fno-zero-initialized-in-bss -ftree-vectorize -fgcse-las
8	h264ref	-O3 -fno-predictive-commoning -fno-function-cse -fsched-rank-heuristic -fmodulo-sched -flto -funswitch-loops -ftree-loop-ivcanon -fsched-dep-count-heuristic -fisolate-erroneous-paths-attribute -fsched-spec-load -fsched-stalled-insns -fcx-limited-range -fno-zero-initialized-in-bss -funroll-all-loops -flive-range-shrinkage -fpeel-loops -fgcse-sm -fweb -fgcse-after-reload -fsched-critical-path-heuristic -fivopts -fvariable-expansion-in-unroller -fno-defer-pop -fipa-pta -fsched-stalled-insns-dep -fno-ira-share-save-slots -fsched-pressure -fbranch-probabilities
9	lbm	-O3 -fmerge-all-constants -floop-unroll-and-jam -fsched-spec-load-dangerous -fvariable-expansion-in-unroller -floop-parallelize-all -fgcse-las -fisolate-erroneous-paths-attribute -fselective-scheduling -fsplit-loops -flimit-function-alignment -fmodulo-sched -fno-ira-share-save-slots -fsched2-use-superblocks -fweb -fno-zero-initialized-in-bss -fpeel-loops -flive-range-shrinkage -fbranch-target-load-optimize2 -fivopts -fgcse-after-reload -fbtr-bb-exclusive -fconserve-stack -fstdarg-opt -fsched-stalled-insns -fno-function-cse -fsection-anchors -freschedule-modulo-scheduled-loops -frename-registers -fdevirtualize-speculatively -fkeep-static-functions -fno-defer-pop -fno-predictive-commoning -fsched-stalled-insns-dep -fsched-last-insn-heuristic -fno-ipa-cp-clone -ftree-vectorize -ftree-loop-ivcanon -fsched-spec-insn-heuristic -fipa-pta -fprefetch-loop-arrays -fprofile-correction -fsched-dep-count-heuristic -fno-ira-share-spill-slots -fsched-rank-heuristic -floop-nest-optimize
10	h263enc	-O3 -flto -fsched-stalled-insns -fcx-limited-range -fgcse-sm -fprefetch-loop-arrays -floop-parallelize-all -fno-ira-share-spill-slots

Flag set	Benchmark	Flags
11	mpeg2dec	-03 -fsched2-use-superblocks -floop-nest-optimize -fsched-pressure -fno-defer-pop -floop-interchange -fno-predictive-commoning -fivopts -fgcse-sm -fsched-stalled-insns-dep -fpeel-loops -fsched-spec-load-dangerous -fsplit-loops -fkeep-static-functions -floop-parallelize-all -fselective-scheduling -fno-ira-share-save-slots -fno-function-cse -funroll-all-loops -fsched-spec-insn-heuristic -flimit-function-alignment -fcx-limited-range -fvpt -fgcse-after-reload -fdevirtualize-speculatively
12	mpeg2enc	-03 -fno-function-cse -fvpt -fsched-group-heuristic -flive-range-shrinkage -fivopts -fsched-stalled-insns -fgcse-las -ftree-parallelize-loops=16 -fgcse-after-reload -funroll-loops -fno-inline -fno-zero-initialized-in-bss -fsched2-use-superblocks -fsched-pressure -fsched-spec-insn-heuristic -fipa-pta -fbranch-target-load-optimize2 -fkeep-inline-functions -fprefetch-loop-arrays -fisolate-erroneous-paths-attribute -ftree-vectorize -fmodulo-sched -funswitch-loops -fdevirtualize-speculatively -ftree-loop-ivcanon -fgcse-sm -fweb -fsched-rank-heuristic -fno-inline-functions
13	overall	-03 -flto -fselective-scheduling -fgcse-sm -floop-unroll-and-jam -flimit-function-alignment -fmodulo-sched -fbranch-target-load-optimize2 -flive-range-shrinkage -fbranch-probabilities -ftree-loop-ivcanon -fno-predictive-commoning -fno-function-cse -fprefetch-loop-arrays -fsched-group-heuristic -floop-parallelize-all -fsched-pressure -fno-zero-initialized-in-bss -fgcse-after-reload -fsched-stalled-insns-dep -fconserve-stack -ftree-parallelize-loops=16 -fsched-critical-path-heuristic -fcx-limited-range -fkeep-inline-functions -fvariable-expansion-in-unroller -frename-registers -fsection-anchors -fbtr-bb-exclusive -fprofile-use -fsplit-loops -fno-ira-share-save-slots -ftree-loop-im -funroll-all-loops -fsched-last-insn-heuristic -fsched-spec-load-dangerous -fkeep-static-functions -fmerge-all-constants -fdevirtualize-speculatively -fweb -fstdarg-opt -fno-ira-share-spill-slots -fbranch-target-load-optimize -floop-interchange
14	consistent	-03 -fselective-scheduling -fsection-anchors -fprefetch-loop-arrays -fdevirtualize-speculatively -fsched-dep-count-heuristic -fweb -fisolate-erroneous-paths-attribute -fpeel-loops -fbranch-target-load-optimize2 -fno-defer-pop -fsched-stalled-insns-dep -fmerge-all-constants -fbtr-bb-exclusive -ftree-loop-ivcanon -fsched-stalled-insns -fno-ira-share-spill-slots -floop-unroll-and-jam -fcx-limited-range -flto -fsched2-use-superblocks -fno-function-cse -frename-registers -fvariable-expansion-in-unroller -fivopts -fgcse-after-reload -floop-interchange -fmodulo-sched -fno-predictive-commoning -fsched-spec-insn-heuristic -fsched-spec-load-dangerous

Table 17: Full flag sets, included in the appendix due to their length